

Robust Algebraic Parameter Estimation via Gaussian Process Regression

Oren Bassik¹ Alexander Demin² Alexey Ovchinnikov^{1,3}

IFAC/COSY/SSSC Joint Conference

July 1, 2025

¹CUNY Graduate Center

²HSE University

³CUNY Queens College

Robust Differentiation: The Cornerstone of Algebraic Parameter Estimation

Challenge: Differentiating Noisy Data

- Algebraic estimators require derivatives $y(t), \dot{y}(t), \ddot{y}(t), \dots$ from noisy measurements $y_k = y(t_k) + \varepsilon_k$
- Naive differentiation acts as a high-pass filter
- Fundamental issue: $\text{Var}[\Delta^n y / \Delta t^n] \propto \sigma^2 \Delta t^{-2n}$
- Result:** Significant noise amplification, degrading derivative quality

*TODO: Plot showing AAA
interpolant
oscillating wildly through noisy data
with corresponding unstable
derivative*

GPR Approach: Probabilistic Differentiation

- Reframe as Bayesian inference: $y(\cdot) \sim \mathcal{GP}(m(\cdot), k(\cdot, \cdot))$
- Smoothness prior via kernel (RBF, Matérn) encodes physical assumptions
- Derivatives computed analytically:
 $\partial^\ell y | \mathcal{D} \sim \mathcal{GP}(\partial^\ell \mu(t), \partial^\ell \partial^{\ell'} \Sigma(t, t'))$
- Result:** Stable derivatives with uncertainty quantification

*TODO: Plot showing smooth GPR
mean
with confidence bands and
corresponding stable derivative*

Key Insight

Evaluating Differentiator Performance

Noise Sensitivity

RMSE vs. Noise Level

TODO: Log-log plot showing:

- *GPR_Julia: stable slope $\alpha \approx 1.1$*
- *AAA: catastrophic slope $\alpha \approx 9.7$ at $\sigma > 10^{-8}$*
 - *95% confidence intervals*
- *Vertical line at failure threshold*

Higher-Order Derivatives

RMSE vs. Derivative Order

TODO: Semi-log plot showing:

- *GPR_Julia: graceful degradation*
- *TVDiff: complete failure at order ≥ 4*
- *Finite differences: exponential growth*
 - *Failure markers for non-convergent cases*

Noise Cliff Analysis

- **GPR_Julia:** Stable performance across all noise levels
- **AAA methods:** High error sensitivity above $\sigma \approx 10^{-8}$
- **Error growth:** GPR error grows linearly; AAA error grows exponentially
- **Statistical sig:** $p < 10^{-12}$ for difference at $\sigma = 10^{-6}$

Derivative Order Analysis

- **GPR_Julia:** Graceful degradation with increasing order
- **TVDiff:** Fails to converge for orders ≥ 4
- **Finite differences:** Rapid error growth

Why Gaussian Process Regression Excels at Numerical Differentiation

Bayesian Framework

1. **Prior over signal:** $f(t) \sim \mathcal{GP}(0, k(t, t'))$

$$\text{RBF kernel: } k(t, t'; \theta) = \sigma_f^2 \exp\left(-\frac{(t-t')^2}{2\ell^2}\right)$$

2. **Posterior inference:** $y = f + \varepsilon, \varepsilon \sim \mathcal{N}(0, \sigma_n^2 I)$

$$\mathbb{E}[f(t^*)|y] = k_*^T (K + \sigma_n^2 I)^{-1} y$$

3. **Analytical derivatives:** $\frac{\partial f}{\partial t} \sim \mathcal{GP}(\cdot, \cdot)$

$$\mathbb{E}[\dot{f}(t^*)|y] = k'_*{}^T (K + \sigma_n^2 I)^{-1} y$$

4. **Hyperparameter optimization:**

$$\log p(y|\theta) = -\frac{1}{2} y^T \Sigma^{-1} y - \frac{1}{2} \log |\Sigma| - \frac{n}{2} \log 2\pi$$

TODO: GPR illustration showing:

- *Noisy data points*
- *GP posterior mean (smooth)*

Key Technical Advantages

1. Probabilistic Function Inference

- Models underlying function $f(t)$, not just data
- Treats differentiation as statistical inference

2. Data-Driven Hyperparameters

- Marginal likelihood automatically finds $\{\sigma_f, \ell, \sigma_n\}$
- No ad-hoc filter tuning required

3. Numerical Stability

- Avoids $\Delta y / \Delta t$ noise amplification
- Regularized kernel matrix: $(K + \sigma_n^2 I)^{-1}$

4. Global Smoothing + Local Accuracy

- All data points inform each estimate
- Kernel decay: exponential locality, matrix coupling: global

Application: Making the Algebraic Method Viable

Benchmarking Against Standard Methods

Test Suite: Nonlinear dynamic systems with realistic noise (1.0% rel. noise)

System	GPR-Algebraic	Original (AAA)	SciML (LM)
Fitzhugh-Nagumo	1.8%	100%	0.5%
Lotka-Volterra	2.1%	100%	0.3%
SEIR Model	1.2%	FAIL	0.4%
HIV Dynamics	3.4%	100%	0.5%

Errors are Mean Relative Error (MRE) in parameter estimates.

SciML is a standard Levenberg-Marquardt optimization solver.

TODO: Trajectory comparison plot

- *Ground truth vs. estimates*
- *GPR-Algebraic (green, accurate)*
- *Original AAA (red, diverging)*
- *Noisy measurements (gray dots)*

Key Contributions

• Enables Robustness

- GPR component overcomes the brittleness of the original algebraic method.

• Preserves Automation

- Retains the key algebraic advantage: no initial parameter guesses required.

• Achieves Viability

- Performance is now competitive with established optimization-based methods.

• Offers Alternate Approach

Conclusions & Outlook

Technical Contributions

- Solved key bottleneck for algebraic method via a principled GPR approach
- Demonstrated noise tolerance
 - GPR-based method is effective up to $\sigma^2 \approx 10^{-2}$
 - Interpolation-based methods fail above $\sigma^2 \approx 10^{-8}$
- Preserved algebraic advantages
 - No initial parameter guesses
 - Fully automated operation
 - Analytic solution guarantees
- Minimal computational overhead
 - $< 20\%$ extra CPU (100ms \rightarrow 118ms)
 - Polynomial solving remains bottleneck
- Identified common failure modes
 - “Noise cliff”: high sensitivity

Scientific Contributions

- Systematic evaluation framework
 - 21 numerical differentiation methods
 - Unified Automatic Differentiation harness
 - Fair, reproducible comparisons
- Identified universal failure modes
 - “Noise cliff”: catastrophic failure at critical σ
 - “Derivative wall”: breakdown at high orders ($k \geq 4$)
- Mathematical foundation
 - Proved GP kernel smoothness \Rightarrow well-posed algebraic estimator
 - Explains unique robustness of GP solution
- Bridges algebraic & Bayesian methods

Future Research Directions

- From Estimates to Confidence
 - Uncertainty quantification via GP posterior variance
 - Credible intervals on parameter estimates
 - Risk-aware parameter estimation
- Data-Driven Experiment Design
 - Optimal experimental design using variance field
 - A- and D-optimality for parameter estimation
 - Active learning for minimal data collection
- Scaling the Framework
 - Sparse/inducing-point GPs $\Rightarrow \mathcal{O}(N \log N)$
 - Extension to 10^2 - 10^3 parameters
 - Real-time implementation

A1: Computational Performance Analysis

Timing Benchmarks

Runtime Comparison (10k samples):

- GPR_Julia: 2.15s
- SavitzkyGolay: 0.0015s
- TVDiff: 0.19s
- FiniteDiff: 0.0006s

Runtime Breakdown (GPR):

- Polynomial solving: 70%
- GPR computation: 20%
- Miscellaneous: 10%

Scaling Analysis

Computational Complexity:

- Naive GP: $\mathcal{O}(N^3)$
- Sparse GP: $\mathcal{O}(NM^2)$, $M \ll N$

*TODO: Bar chart showing
runtime breakdown and
method comparison*

Key Insight

GPR is **not** the computational bottleneck.

Accuracy-first approach justified by
minimal overhead vs. dramatic robustness gains.

A2: Extended Benchmark Results

Complete Performance Matrix: 21 Methods \times 3 Noise Levels \times 6 Derivative Orders

*TODO: Comprehensive heatmap showing
 $\log_{10}(\text{RMSE})$ across all conditions
Red = failure, Green = success
Generated by method_comparison_heatmap.pdf*

Statistical Analysis

- **Failure Rate Analysis:** GPR-Julia: 0% (0/63), AAA-Julia: 37% (23/63), TVDiff: 52% (33/63)
- **Edge Cases:** Very low noise ($\sigma < 10^{-10}$) requires numerical conditioning
- **Confidence Intervals:** 95% CI on geometric mean RMSE across 30 MC runs per condition
- **Method Categories:** Finite differences, spectral, variational, kernel-based

A3: GPR Implementation & Methodology Details

Hyperparameter Optimization

Marginal Likelihood Maximization:

$$\log p(y|\theta) = -\frac{1}{2}y^T \Sigma^{-1}y - \frac{1}{2} \log |\Sigma| - \frac{n}{2} \log 2\pi$$

Optimization Strategy:

- L-BFGS-B with multiple random initializations
- Parameter bounds: $\ell \in [10^{-3}, 10^3]$, $\sigma_f^2 \in [10^{-6}, 10^6]$
- Automatic noise floor: $\sigma_n^2 \geq 10^{-12}$

Kernel Selection

RBF vs. Matérn Comparison:

- RBF: C^∞ smooth, best for high-order derivatives
- Matérn-3/2: C^1 smooth, computational efficiency
- Matérn-5/2: C^2 smooth, good compromise

Integration Framework

Automatic Differentiation:

- Julia: ForwardDiff.jl for exact derivatives
- Python: JAX for vectorized operations
- Enables fair, consistent method comparison

Software Integration:

- Modular design: derivative estimator + algebraic solver
- Compatible with existing parameter estimation pipelines
- Export to MATLAB, Python, Julia formats

TODO: Workflow diagram
Data \rightarrow GPR \rightarrow Derivatives \rightarrow
Algebraic Solver \rightarrow Parameters